

VGF: Value-Guided Fuzzing

Fuzzing Hardware as Hardware

Ruochen Dai[†], Michael Lee[†], Patrick Hoey[‡], Weimin Fu^{*}, Yier Jin[¶], Xiaolong Guo^{*}, Shuo Wang[†], Dean Sullivan[§], Tuba Yavuz[†], Orlando Arias[‡]

[†] University of Florida

[‡] University of Massachusetts Lowell

^{*} Kansas State University

[¶] University of Science and Technology of China

[§] University of New Hampshire

Have less buggy hardware

Have less buggy hardware

Dynamically generate test vectors to check assertions.

The Issue With Static Analysis

Explosion of states.


The Issue With Static Analysis

Explosion of states.




The Issue With Static Analysis

Explosion of states.




The Issue With Static Analysis

Explosion of states.




The Issue With Static Analysis

Explosion of states.




The Issue With Static Analysis

Explosion of states.




The Issue With Static Analysis

Explosion of states.




Dynamic Analysis of Hardware?




???


Dynamic Analysis of Hardware?



Dynamic Analysis of Hardware?



Dynamic Analysis of Hardware?



Fuzz Testing Hardware?

Previous work

- ▶ Convert hardware model into C/C++ model
- ▶ Instrument resulting code
- ▶ Use something like AFL to fuzz

Fuzz Testing Hardware?

Previous work

- ▶ Convert hardware model into C/C++ model
- ▶ Instrument resulting code
- ▶ Use something like AFL to fuzz

Survey Paper



W. Fu, O. Arias, Y. Jin, and X. Guo, "Fuzzing hardware: Faith or reality? : Invited paper," in *2021 IEEE/ACM International Symposium on Nanoscale Architectures (NANOARCH)*, 2021, pp. 1–6

A Closer Look at Previous Work

Assume that conversion into C/C++ model is 100% accurate.

Use basic block as measure of coverage.

Use lines visited as measure of coverage.

Does the approach always work?

A Closer Look at Previous Work

Assume that conversion into C/C++ model is 100% accurate.

Use basic block as measure of coverage.

Use lines visited as measure of coverage.

Does the approach always work?

Implicit control-flow

```
module mux(output o,  
           input a, b, s);  
    assign o = s ? b : a;  
endmodule
```

No control-flow

```
module mux(output o,  
           input a, b, s);  
    assign o = |{a & ~s, b & s};  
endmodule
```

The Many Faces of a Multiplexer

```
module mux(output o,  
    input a, b, c);  
    assign o = (a & ~s) | (b & s);  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    assign o = |{a & ~s, b & s};  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    wire t[1:0];  
    and u0(t[0], a, ~s);  
    and u1(t[1], b, s);  
    or u2(o, t[0], t[1]);  
endmodule
```

```
primitive mux(o, a, b, s);  
    output o;  
    input a, b, s;  
    table  
        0 ? 0 : 0;  
        1 ? 0 : 1;  
        ? 0 1 : 0;  
        ? 1 1 : 1;  
        ? ? x : x;  
    endtable  
endprimitive
```

```
module mux(output logic o,  
    input a, b, c);  
    always_comb @(*)  
        if(s) o = b;  
        else o = a;  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    assign o = s ? b : a;  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    wire s_inv;  
    pmos u0(s_inv, 1, s);  
    nmos u1(s_inv, 0, s);  
    cmos u2(o, a, s_inv, s);  
    cmos u3(o, b, s, s_inv);  
endmodule
```

The Many Faces of a Multiplexer

```
module mux(output o,  
    input a, b, c);  
    assign o = (a & ~s) | (b & s);  
endmodule
```

```
module mux(output logic o,  
    input a, b, c);  
    always_comb @(*)  
        if(s) o = b;  
        else o = a;  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    assign o = |{a & ~s, b & s};  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    assign o = s ? b : a;  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    wire t[1:0];  
    and u0(t[0], a, ~s);  
    and u1(t[1], b, s);  
    or u2(o, t[0], t[1]);  
endmodule
```

```
primitive mux(o, a, b, s);  
    output o;  
    input a, b, s;  
    table  
        0 ? 0 : 0;  
        1 ? 0 : 1;  
        ? 0 1 : 0;  
        ? 1 1 : 1;  
        ? ? x : x;  
    endtable  
endprimitive
```

```
module mux(output o,  
    input a, b, c);  
    wire s_inv;  
    pmos u0(s_inv, 1, s);  
    nmos u1(s_inv, 0, s);  
    cmos u2(o, a, s_inv, s);  
    cmos u3(o, b, s, s_inv);  
endmodule
```

The Many Faces of a Multiplexer

```
module mux(output o,  
    input a, b, c);  
    assign o = (a & ~s) | (b & s);  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    assign o = |{a & ~s, b & s};  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    wire t[1:0];  
    and u0(t[0], a, ~s);  
    and u1(t[1], b, s);  
    or u2(o, t[0], t[1]);  
endmodule
```


```
primitive mux(o, a, b, s);  
    output o;  
    input a, b, s;  
    table  
        0 ? 0 : 0;  
        1 ? 0 : 1;  
        ? 0 1 : 0;  
        ? 1 1 : 1;  
        ? ? x : x;  
    endtable  
endprimitive
```

```
module mux(output logic o,  
    input a, b, c);  
    always_comb @(*)  
        if(s) o = b;  
        else o = a;  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    assign o = s ? b : a;  
endmodule
```

```
module mux(output o,  
    input a, b, c);  
    wire s_inv;  
    pmos u0(s_inv, 1, s);  
    nmos u1(s_inv, 0, s);  
    cmos u2(o, a, s_inv, s);  
    cmos u3(o, b, s, s_inv);  
endmodule
```

Design Coverage?




Gate Level Circuit

```
module mux(output o,  
    input a, b, s);  
    wire t[1:0];  
    and u0(t[0], a, ~s);  
    and u1(t[1], b, s);  
    or u2(o, t[0], t[1]);  
endmodule
```

Equivalent High-Level Circuit

```
module mux(output logic o,  
    input a, b, s);  
    always_comb @(*)  
        if(s) o = b;  
        else o = a;  
endmodule
```

Design Coverage?




Gate Level Circuit

```
module mux(output o,  
    input a, b, s);  
    wire t[1:0];  
    and u0(t[0], a, ~s);  
    and u1(t[1], b, s);  
    or u2(o, t[0], t[1]);  
endmodule
```

Equivalent High-Level Circuit

```
module mux(output logic o,  
    input a, b, s);  
    always_comb @(*)  
        if(s) o = b;  
        else o = a;  
endmodule
```

Design Coverage?




Gate Level Circuit

```
module mux(output o,  
    input a, b, s);  
    wire t[1:0];  
    and u0(t[0], a, ~s);  
    and u1(t[1], b, s);  
    or u2(o, t[0], t[1]);  
endmodule
```

Equivalent High-Level Circuit

```
module mux(output logic o,  
    input a, b, s);  
    always_comb @(*)  
        if(s) o = b;  
        else o = a;  
endmodule
```


Design Coverage?




Gate Level Circuit

```
module mux(output o,
            input a, b, s);
    wire t[1:0];
    and u0(t[0], a, ~s);
    and u1(t[1], b, s);
    or u2(o, t[0], t[1]);
endmodule
```

Equivalent High-Level Circuit

```
module mux(output logic o,
            input a, b, s);
    always_comb @(*)
        if(s) o = b;
        else o = a;
endmodule
```

Design Coverage?



Gate Level Circuit

```
module mux(output o,  
    input a, b, s);  
    wire t[1:0];  
    and u0(t[0], a, ~s);  
    and u1(t[1], b, s);  
    or u2(o, t[0], t[1]);  
endmodule
```

Equivalent High-Level Circuit

```
module mux(output logic o,  
    input a, b, s);  
    always_comb @(*)  
        if(s) o = b;  
        else o = a;  
endmodule
```

Looking Closer at Coverage

Gate Level Circuit

```
module mux(output o,  
           input a, b, s);  
  wire t[1:0];  
  and u0(t[0], a, ~s);  
  and u1(t[1], b, s);  
  or u2(o, t[0], t[1]);  
endmodule
```

Equivalent High-Level Circuit

```
module mux(output logic o,  
           input a, b, s);  
  always_comb @(*)  
    if(s) o = b;  
    else o = a;  
endmodule
```

Observations

- ▶ Gate level circuit has 100% line coverage on a single simulation tick
- ▶ Gate level circuit has no control-flow information

Verilated Output?

Looking at the C++ conversion of each HDL style:

Gate Level Circuit

```
vlSelfRef.o = (((~ (IData)(vlSelfRef.s)) & (IData)(vlSelfRef.a))  
               | ((IData)(vlSelfRef.b) & (IData)(vlSelfRef.s)));
```

Equivalent High-Level Circuit

```
vlSelfRef.o = ((IData)(vlSelfRef.s)  
               ?  
                 (IData)(vlSelfRef.b)  
               :  
                 (IData)(vlSelfRef.a));
```

The Insight

Instead of looking at lines of code or control-flow information, look at the state space of the design.

Tracking Signal Values

Can we obtain the values of signals at simulation time?

Tracking Signal Values

Can we obtain the values of signals at simulation time?

- ▶ **Yes!**
- ▶ Most HDL simulators implement the Verilog Procedural Interface
- ▶ Use VPI to interface with and extend simulator!

Tracking Signal Values

Can we obtain the values of signals at simulation time?

- ▶ **Yes!**
- ▶ Most HDL simulators implement the Verilog Procedural Interface
- ▶ Use VPI to interface with and extend simulator!
- ▶ Have a design-agnostic VPI harness, launch simulator with harness

Tracking Signal Values

Can we obtain the values of signals at simulation time?

- ▶ **Yes!**
- ▶ Most HDL simulators implement the Verilog Procedural Interface
- ▶ Use VPI to interface with and extend simulator!
- ▶ Have a design-agnostic VPI harness, launch simulator with harness
- ▶ Integrate harness with a fuzzer for input generation

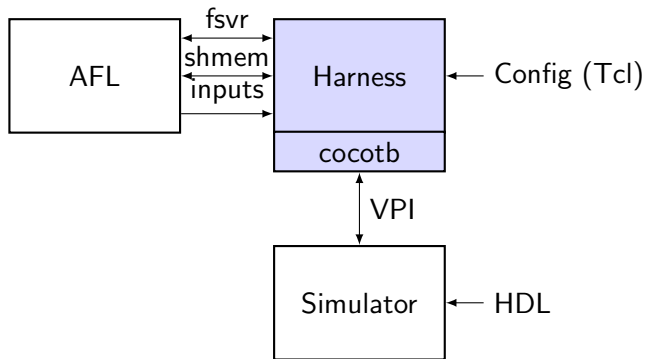
VGF Methodology

- ▶ Tell harness what assertions to check for
- ▶ Tell harness which signals to track
- ▶ Tie harness to a fuzzing engine
- ▶ Use signal changes to provide feedback to the fuzzer

VGF Methodology

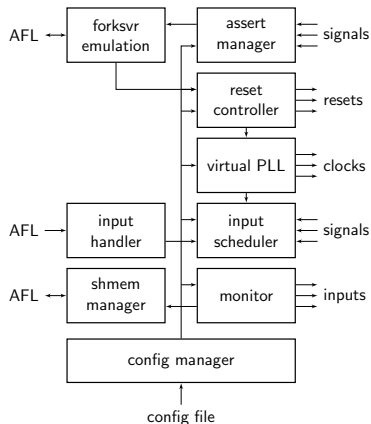
- ▶ Tell harness what assertions to check for
- ▶ Tell harness which signals to track
- ▶ Tie harness to a fuzzing engine
- ▶ Use signal changes to provide feedback to the fuzzer
- ▶ Let the system run

VGF Architecture



Harness Architecture

- ▶ Use VPI to communicate with DUT
- ▶ Communication with AFL done through UNIX pipes and shared memory
- ▶ Harness configuration done using Tcl script



The AFL Instrumentation

```
if(condition) {
    current_location = <COMPILER_RANDOM_NUMBER_0>;
    shared_memory[current_location ^ previous_location]++;
    previous_location = current_location >> 1;
    /* ... */
} else {
    current_location = <COMPILER_RANDOM_NUMBER_1>;
    shared_memory[current_location ^ previous_location]++;
    previous_location = current_location >> 1;
    /* ... */
}
```

The AFL Instrumentation

- ▶ 64 KiB allocation as `shared_memory`
 - ▶ Name exported in environment variable `__AFL_SHM_ID`
- ▶ Shift done to avoid issues with loops:

```
do {  
    current_location = <COMPILER_RANDOM_NUMBER_2>;  
    shared_memory[current_location ^ previous_location]++;  
    previous_location = current_location >> 1;  
    /* code with no branches or loops */  
} while(condition);
```

if not using shift, `0 == current_location ^ previous_location`

AFL Deciding on Input

For every bucket in the control-flow hash map, AFL keeps a *preferred* input queue. The input queue contains a copy of the data fed as input to the fuzzed program.

\emptyset	q_3	q_1	q_1	\emptyset	\dots	\emptyset	q_2
-------------	-------	-------	-------	-------------	---------	-------------	-------

If a more favorable input is found for a particular bucket, AFL replaces the queue for that bucket with the more favorable one.

AFL Reward Function

- ▶ AFL prioritizes inputs that are small and yield fast executions
- ▶ For a given input queue q , AFL computes

$\text{fav_score} = q \rightarrow \text{exec_us} * q \rightarrow \text{len}$

and tries to minimize this number.

- ▶ Number of buckets is not used on this reward

AFL Reward Function

- ▶ AFL prioritizes inputs that are small and yield fast executions
- ▶ For a given input queue q , AFL computes

$\text{fav_score} = q \rightarrow \text{exec_us} * q \rightarrow \text{len}$

and tries to minimize this number.

- ▶ Number of buckets is not used on this reward
- ▶ More triggers that need to be resolved by the simulator imply that execution is slower!
 - ▶ Recall that we are using simulators that actually obey timing things like #10 **assign** $a = b$;

AFL Reward Function

- ▶ AFL prioritizes inputs that are small and yield fast executions
- ▶ For a given input queue q , AFL computes

$\text{fav_score} = q \rightarrow \text{exec_us} * q \rightarrow \text{len}$

and tries to minimize this number.

- ▶ Number of buckets is not used on this reward
- ▶ More triggers that need to be resolved by the simulator imply that execution is slower!
 - ▶ Recall that we are using simulators that actually obey timing things like `#10 assign a = b;`
- ▶ Reward function is unsuitable for our purposes!

New Reward Functions

New reward functions have been added to AFL.

- ▶ Geometric mean: $\sqrt[n]{\prod^n b_k}$
- ▶ Arithmetic mean: $\frac{1}{n} \sum^n b_k$
- ▶ Total value: $\sum^n b_k$
- ▶ Total buckets: n

where n is the number of non-zero buckets in the hash map and b_k the value in a non-zero bucket. We try to maximize these values.

VGF vs. Formal Verification


Design	VGF	EBMC	FPV
AES128-T100	5.2	— [†]	20.09
lock_case	11.6	0.11	18.41
lock_micro	9.8	422.29	17.16
DAIO	3.8	1.03	17.37
Dekker	4	0.18	17.25
Unidec	— [†]	— [‡]	17.92
AES128-T2500	4.2	— [†]	25.04
RS232-T600	4	0.31	17.35
AE18_core	5	10.66	17.11
async_fifo	93.8	— [‡]	17.37

[†] Allotted time limit reached, timeout

[‡] Tool encountered a parsing error

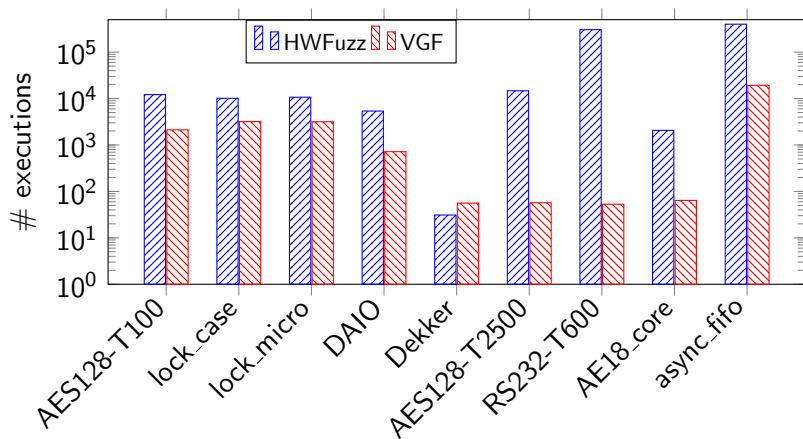
Table: Time to assertion trigger (in s) of VGF compared to Enhanced Bounded Model Checker (EBMC), and Synopsys VC Formal FPV.

VGF vs. Hardware Fuzzing



T. Trippel, K. G. Shin, A. Chernyakhovsky, G. Kelly, D. Rizzo, and M. Hicks, "Fuzzing hardware like software," in *31st USENIX Security Symposium (USENIX Security 22)*, 2022, pp. 3237–3254

VGF vs. Hardware Fuzzing



T. Trippel, K. G. Shin, A. Chernyakhovsky, G. Kelly, D. Rizzo, and M. Hicks, “Fuzzing hardware like software,” in *31st USENIX Security Symposium (USENIX Security 22)*, 2022, pp. 3237–3254

Conclusions

- ▶ It is possible to build a design agnostic *and* language agnostic simulation harness
- ▶ It is possible to trap hardware bugs with dynamic test vector generation
- ▶ We are faster than both open source fuzzers as well as formal verification tools

Conclusions

- ▶ It is possible to build a design agnostic *and* language agnostic simulation harness
- ▶ It is possible to trap hardware bugs with dynamic test vector generation
- ▶ We are faster than both open source fuzzers as well as formal verification tools
- ▶ More testing is needed, larger designs are welcomed

Questions?

orlando_arias@uml.edu